

```
[1]: #Rachit Manandhar  
#2501387  
#np03cs4a240053
```

```
[2]: import pandas as pd  
import numpy as np
```

Problem 1 - Perform a classification task with knn from scratch

Load the Dataset:

```
[6]: # Loading the dataset  
df = pd.read_csv("diabetes.csv")  
  
# Understanding the dataset  
# Show first rows  
df.head()
```

```
[6]:   Pregnancies Glucose BloodPressure SkinThickness Insulin BMI DiabetesPedigreeFunction Age Outcome  
0         6     148           72          35      0  33.6           0.627    50       1  
1         1      85            66          29      0  26.6           0.351    31       0  
2         8     183           64            0      0  23.3           0.672    32       1  
3         1      89            66          23     94  28.1           0.167    21       0  
4         0     137           40          35     168  43.1           2.288    33       1
```

Performing EDA

```
[7]: # Checking dataset info  
df.info()  
  
# Summary  
df.describe()  
  
# Checking Missing values  
df.isna().sum()  
  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 768 entries, 0 to 767  
Data columns (total 9 columns):  
 #   Column           Non-Null Count  Dtype     
---    
 0   Pregnancies      768 non-null    int64    
 1   Glucose          768 non-null    int64    
 2   BloodPressure    768 non-null    int64    
 3   SkinThickness    768 non-null    int64    
 4   Insulin          768 non-null    int64    
 5   BMI              768 non-null    float64  
 6   DiabetesPedigreeFunction 768 non-null    float64  
 7   Age              768 non-null    int64    
 8   Outcome          768 non-null    int64    
dtypes: float64(2), int64(7)  
memory usage: 54.1 KB  
  
[7]: Pregnancies      0  
Glucose          0  
BloodPressure    0  
SkinThickness    0  
Insulin          0  
BMI              0  
DiabetesPedigreeFunction 0  
Age              0  
Outcome          0  
dtype: int64
```

Handle Missing data:

```
[9]: # Replacing missing data since there are not any null values but some data are encoded as 0 which is impossible  
  
# Taking potential columns where 0 is not possible  
col_with_zeros = ["Glucose", "BloodPressure", "SkinThickness", "Insulin", "BMI"]  
  
df_clean = df.copy()  
  
#Replacing the 0s with null first then filling them with median using .fillna()  
for col in col_with_zeros:  
    df_clean[col] = df_clean[col].replace(0, np.nan)  
    df_clean[col] = df_clean[col].fillna(df_clean[col].median())
```

Feature Engineering

```
[30]: feature_matrix_X = df_clean.drop(columns=["Outcome"]).values
target_Y = df_clean["Outcome"].values

# Shuffling the indexes for selecting random data for training and testing
indices = np.arange(len(feature_matrix_X))
np.random.shuffle(indices)

# Splitting the dataset for training and testing
split_size = int(0.7 * len(feature_matrix_X))

train_indices = indices[:split_size]
test_indices = indices[split_size:]

x_train, x_test = feature_matrix_X[train_indices], feature_matrix_X[test_indices]
y_train, y_test = target_Y[train_indices], target_Y[test_indices]
```

Implement kNN:

Distance function using Euclidean distance

```
[32]: def euclidean_distance(a, b):
    """
    Calculates the Euclidean distance between two points in n-dimensional space.

    Arguments:
    a: np.ndarray
        The first point as a numpy array.
    b: np.ndarray
        The second point as a numpy array.

    Returns:
    float
        The Euclidean distance between the two points.
    """

    # Calculating the Euclidean distance
    distance = np.sqrt(np.sum((a - b)**2, axis=1))
    return distance
```

Predicting the class for a single query

```
[15]: def predict_single(x, x_train, y_train, k = 3):
    """
    Predicts the class label of a single query using k-NN algorithm.

    Arguments:
    x: np.ndarray
        The query point for which the prediction is to be made.
    x_train: np.ndarray
        The training feature matrix.
    y_train: np.ndarray
        The training labels.
    k: int
        The number of nearest neighbors to consider (default = 3).

    Returns:
    int
        The predicted class label for the query.
    """

    # Get all the distances to all the training samples
    distances = euclidean_distance(x_train, x)

    sorted_indices = distances.argsort()[:k]
    sorted_labels = y_train[sorted_indices]

    prediction = np.bincount(sorted_labels).argmax()
    return prediction
```

Predicting classes for all test samples

```
[18]: def predict_all(x_test, x_train, y_train, k = 3):
    """
    Predicts the class labels for all test samples using k-NN algorithm.

    Arguments:
    x_test: np.ndarray
        The test feature matrix.
    x_train: np.ndarray
        The training feature matrix.
    y_train: np.ndarray
        The training labels.
    k: int
        The number of nearest neighbors to consider (default = 3).

    Returns:
    np.ndarray
        An array of predicted class labels for the test samples.
    """

    predictions = [predict_single(x, x_train, y_train, k) for x in x_test]
    return np.array(predictions)
```

Evaluating performance using accuracy

```
[22]: def calculate_accuracy(y_true, y_pred):
    """
    Calculates the accuracy of predictions.

    Arguments:
    y_true: np.ndarray
        The true labels.
    y_pred: np.ndarray
        The predicted labels.

    Returns:
    float:
        The accuracy as percentage.
    """

    correct_predictions = np.sum(y_true == y_pred)
    total = len(y_true)
    accuracy = (correct_predictions / total) * 100
    return accuracy
```

Testing

```
[46]: pred = predict_all(x_test, x_train, y_train, 3)
accuracy = calculate_accuracy(y_test, pred)
print(f"Accuracy= {accuracy:.2f}%")
```

Accuracy= 74.89%

Problem 2: Experimentation:

Scaling

```
[34]: # Scaling
means = x_train.mean(axis = 0)
stds = x_train.std(axis = 0)
stds[stds == 0] = 1 # Avoiding division by 0.

# Scaled data training and testing sets
x_train_scaled = (x_train - means) / stds
x_test_scaled = (x_test - means) / stds
```

Using Scaled Data:

```
[45]: scaled_pred = predict_all(x_test_scaled, x_train_scaled, y_train, 3)
scaled_accuracy = calculate_accuracy(y_test, scaled_pred)
print(f"Scaled accuracy= {scaled_accuracy:.2f}%")
Scaled accuracy= 80.09%
```

Comparative Analysis:

Accuracy without scaling : 74.89%

Accuracy with scaling: 80.09%

Scaling impacted the k-NN performance by balancing and giving more classification.

This is because k-NN is a distance based algorithm which means if there is a feature with larger numeric range then it dominates the distance even if it might not be important.

Problem 3: Experimentation with k:

```
[54]: import time
import matplotlib.pyplot as plt
```

Varying the number of neighbors - k:

```
[52]: original_accuracies = []
scaled_accuracies = []

original_time = []
scaled_time = []

k_values = range(1,16)

for k in k_values:
    # For original dataset
    start_time = time.time()
    original_pred = predict_all(x_test, x_train, y_train, k)
    end_time = time.time()

    original_accuracy = calculate_accuracy(y_test, original_pred)
    original_accuracies.append(original_accuracy)
    original_time.append(end_time - start_time)

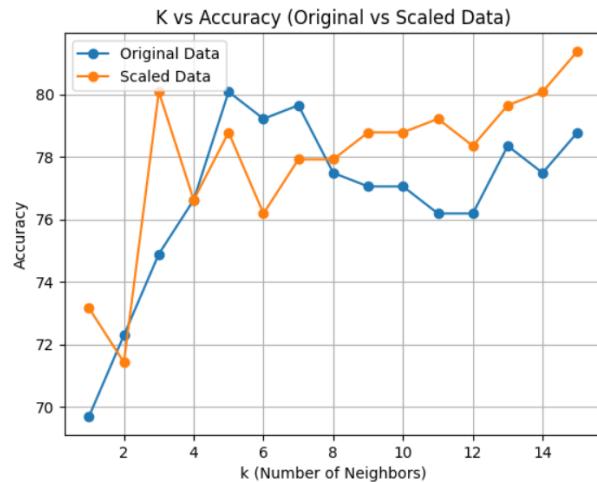
    # For scaled dataset
    start_time = time.time()
    scaled_pred = predict_all(x_test_scaled, x_train_scaled, y_train, k)
    end_time = time.time()

    scaled_accuracy = calculate_accuracy(y_test, scaled_pred)
    scaled_accuracies.append(scaled_accuracy)
    scaled_time.append(end_time - start_time)
```

Visualizing the Results:

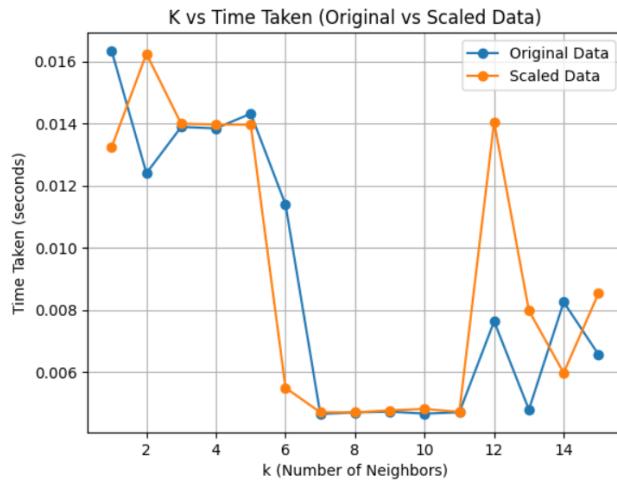
```
[55]: plt.plot(k_values, original_accuracies, marker='o', label='Original Data')
plt.plot(k_values, scaled_accuracies, marker='o', label='Scaled Data')

plt.xlabel("k (Number of Neighbors)")
plt.ylabel("Accuracy")
plt.title("K vs Accuracy (Original vs Scaled Data)")
plt.grid(True)
plt.legend()
plt.show()
```



```
[57]: plt.plot(k_values, original_time, marker='o', label='Original Data')
plt.plot(k_values, scaled_time, marker='o', label='Scaled Data')

plt.xlabel("k (Number of Neighbors)")
plt.ylabel("Time Taken (seconds)")
plt.title("K vs Time Taken (Original vs Scaled Data)")
plt.grid(True)
plt.legend()
plt.show()
```



Analysis

The choice of k affects the accuracy as small k value makes the model overfitting and sensitive to noise while large k value makes the model underfitting and drops in accuracy. So, it is crucial for selecting optimal k value which often gives the best accuracy. The choice of k does not significantly affect computational cost as k is used only for selecting/sorting while calculating the distance takes the most computational cost.

Based on my analysis: The optimal value of k is 5.

Problem 4 - Additional Questions

Challenges of using k-NN for large datasets and high-dimensional data are;

High computation cost: For every test sample, it calculates the distance between the test sample and all the training samples and sort them. High Memory Usages: k-NN stores the entire training dataset in the memory rather than learning a model which means for large datasets the memory requirement would also be high. Dimensionality: In high dimensional data, the distance between the points become meaningless and unreliable. Sensitive to scaling: Features with large numeric ranges dominates the distance calculation even if its not important.