

```
# Rachit Manandhar
```

```
# 2501387
```

```
# np03cs4a240053
```

```
[14]: # Rachit Manandhar  
# 2501387  
# np03cs4a240053
```

```
[4]: import pandas as pd  
import numpy as np
```

### 3.1.1 Data Understanding, Analysis and Preparations:

To - Do - 1:

```
[5]: # 1  
df = pd.read_csv("student.csv")
```

```
[21]: # 2  
print("First 5 rows: ")  
print(df.head())  
print("\nLast 5 rows: ")  
print(df.tail())
```

```
First 5 rows:  
   Math  Reading  Writing  
0     48       68      63  
1     62       81      72  
2     79       80      78  
3     76       83      79  
4     59       64      62
```

```
Last 5 rows:  
   Math  Reading  Writing  
995    72       74      70  
996    73       86      90  
997    89       87      94  
998    83       82      78  
999    66       66      72
```

```
[22]: # 3  
print("Info: ")  
df.info()
```

```
Info:  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1000 entries, 0 to 999  
Data columns (total 3 columns):  
 #   Column  Non-Null Count  Dtype    
---  --    
 0   Math    1000 non-null  int64  
 1   Reading 1000 non-null  int64  
 2   Writing 1000 non-null  int64  
dtypes: int64(3)  
memory usage: 23.6 KB
```

```
[23]: # 4  
print("Descriptive info: ")  
df.describe()
```

Descriptive info:

	Math	Reading	Writing
count	1000.000000	1000.000000	1000.000000
mean	67.290000	69.872000	68.616000
std	15.085008	14.657027	15.241287
min	13.000000	19.000000	14.000000
25%	58.000000	60.750000	58.000000
50%	68.000000	70.000000	69.500000
75%	78.000000	81.000000	79.000000
max	100.000000	100.000000	100.000000

```
[6]: # 5  
features_X = df.drop(columns = ["Writing"]).values  
label_Y = df["Writing"].values
```

## To - Do - 2:

```
[7]: X = features_X.T  
Y = label_Y  
  
d = X.shape[0]  
W = np.zeros((d, 1))  
  
y_pred = (W.T @ X).T  
  
print("X shape:", X.shape)      # (d, n)  
print("W shape:", W.shape)      # (d, 1)  
print("Y shape:", Y.shape)      # (n, 1)  
print("Y_pred shape:", y_pred.shape)
```

```
X shape: (2, 1000)  
W shape: (2, 1)  
Y shape: (1000,)  
Y_pred shape: (1000, 1)
```

## To - Do - 3:

```
[9]: def train_test_split(X, Y, test_size = 0.2, random_state=42):  
    indices = np.arange(len(X))  
    np.random.seed(random_state)  
    np.random.shuffle(indices)  
  
    split = int(test_size * len(X))  
  
    train_indices = indices[split:]  
    test_indices = indices[:split]  
  
    return X[train_indices], X[test_indices], Y[train_indices], Y[test_indices]  
  
x_train, x_test, y_train, y_test = train_test_split(features_X, label_Y, test_size = 0.2, random_state = 42)
```

## To - Do - 4:

```
[10]: def cost_function(X, Y, W):
    """
    Calculates the Mean Square Error (MSE)

    Arguments:
    X: array-like, shape (n_samples, n_features)
        Feature Maxtrix
    Y: array-like, shape (n_samples,)
        True target values.
    W: array-like, shape (n_features, )
        Weight vector

    Returns:
    float
        The Mean Squared Error (MSE)
    """
    X = np.array(X, dtype = float)
    Y = np.array(Y, dtype = float).reshape(-1, 1)
    W = np.array(W, dtype = float).reshape(-1, 1)

    n = len(Y)

    Y_predicted = X @ W
    error = Y_predicted - Y
    MSE = (1 / (2 * n)) * np.sum(error ** 2)

    return MSE
```

## To - Do - 5:

```
[11]: # Test case
X_test = np.array([[1, 2], [3, 4], [5, 6]])
Y_test = np.array([3, 7, 11])
W_test = np.array([1, 1])
cost = cost_function(X_test, Y_test, W_test)
if cost == 0:
    print("Proceed Further")
else:
    print("something went wrong: Reimplement a cost function")
print("Cost function output:", cost_function(X_test, Y_test, W_test))
```

Proceed Further

Cost function output: 0.0

## To - Do - 6:

```
[12]: def gradient_descent(X, Y, W, alpha, iterations):
    """
    Perform gradient descent to optimize the parameters of a linear regression model.

    Parameters:
        X (numpy.ndarray): Feature matrix (m x n).
        Y (numpy.ndarray): Target vector (m x 1).
        W (numpy.ndarray): Initial guess for parameters (n x 1).
        alpha (float): Learning rate.
        iterations (int): Number of iterations for gradient descent.

    Returns:
        tuple: A tuple containing the final optimized parameters (W_update) and the history of cost values.
        W_update (numpy.ndarray): Updated parameters (n x 1).
        cost_history (list): History of cost values over iterations.
    """
    X = np.array(X,dtype=float)
    Y= np.array(Y,dtype=float).reshape(-1,1)
    W= np.array(W,dtype=float).reshape(-1,1)

    m= len(Y)
    cost_history = [] # To store cost at each iteration
    W_update = W.copy()

    for iteration in range(iterations):
        # Step 1: Hypothesis values
        Y_pred = X @ W_update

        # Step 2: Difference between hypothesis and actual Y
        loss = Y_pred - Y

        # Step 3: Gradient calculation
        dw = (1/m) * (X.T @ loss)

        # Step 4: Update W
        W_update = W_update - alpha * dw

        # Step 5: Compute new cost
        cost = cost_function(X, Y, W_update)
        cost_history.append(cost)

        # # PRINT one line per iteration
        # print(f"Iteration {iteration+1}:")
        # print(" Weights:\n", W_update)
        # print(" Cost:", cost)
        # print("-" * 30)

    return W_update, cost_history
```

## To - Do - 7:

```
[13]: # Generate random test data
np.random.seed(0) # For reproducibility
X = np.random.rand(100, 3) # 100 samples, 3 features
Y = np.random.rand(100)
W = np.random.rand(3) # Initial guess for parameters
# Set hyperparameters
alpha = 0.01
iterations = 1000
# Test the gradient_descent function
final_params, cost_history = gradient_descent(X, Y, W, alpha, iterations)
# Print the final parameters and cost history
print("Final Parameters:", final_params)
print("Cost History:", cost_history)

Final Parameters: [[0.20551667]
 [0.54295081]
 [0.10388027]]
Cost History: [np.float64(0.10711197094660153), np.float64(0.10634880599939901), np.float64(0.10559826315680618), np.float64(0.10486012948320558), np.float64(0.1041341956428534), np.float64(0.1034202558390062), np.float64(0.1027181077540776), np.float64(0.1020275524908062), np.float64(0.1013483945144193), np.float64(0.1006804415957737), np.float64(0.1000235047554587), np.float64(0.09937739820884377), np.float64(0.09874193931205609), np.float64(0.09811694850887098), np.float64(0.09750224927850094), np.float64(0.0968976680842672), np.float64(0.09630303432313951), np.float64(0.09571818027612913), np.float64(0.09514294105952065), np.float64(0.09457715457692842), np.float64(0.09402066147216397), np.float64(0.0934730508290017), np.float64(0.0929349313951191), np.float64(0.09240538899833017), np.float64(0.09188452904154543), np.float64(0.0913722051899995), np.float64(0.090868627358260123), np.float64(0.0903259279010509), np.float64(0.0898850237398919), np.float64(0.08940542984603007), np.float64(0.08893367662855953), np.float64(0.08846963201539432), np.float64(0.08801316613342668), np.float64(0.0875645130486386), np.float64(0.08712246201010665), np.float64(0.08668797485125508), np.float64(0.08626056851623207), np.float64(0.08584012374351278), np.float64(0.08542652328745133), np.float64(0.08501965188419301), np.float64(0.0846193962181636), np.float64(0.0844488912489), np.float64(0.08383828837978763), np.float64(0.08345721902397185), np.float64(0.08308233097530582), np.float64(0.08271352017645425), np.float64(0.08235068432886682), np.float64(0.0819372286303817), np.float64(0.08164253690927113), np.float64(0.08129702926893387), np.float64(0.08095710438620353), np.float64(0.08062266832028739), np.float64(0.08029362871811391), np.float64(0.07996989478748553), np.float64(0.0796513772706855), np.float64(0.07933798841853089), np.float64(0.07902964196486459), np.float64(0.07872625310147845), np.float64(0.07842773845346054), np.float64(0.078050053253578), np.float64(0.0775666270458499), np.float64(0.077280803333641404), np.float64(0.07700545763317514), np.float64(0.07673451466614989), np.float64(0.07646790043736812), np.float64(0.07620554219936448), np.float64(0.07594736843403344), np.float64(0.07569330883184205), np.float64(0.07544329427139428), np.float64(0.07519725679934074), np.float64(0.07495512961062821), np.float64(0.07471684702980327), np.float64(0.0744823448832412), np.float64(0.0742515585129952), np.float64(0.07402442670031911), np.float64(0.0738008877019607), np.float64(0.07358088120619749), np.float64(0.0733643479203919), np.float64(0.07315122955375959), np.float64(0.07294146880042966), np.float64(0.07273500932279067), np.float64(0.07253179573511871), np.float64(0.07233177358748233), np.float64(0.0721348893499193), np.float64(0.07194109936688139), np.float64(0.0717503249919482), np.float64(0.07156254227276149), np.float64(0.07137769223630935), np.float64(0.07119572572433286), np.float64(0.0710594209097385), np.float64(0.07084025077922623), np.float64(0.07066648126131), np.float64(0.07049574053020462), np.float64(0.07032748284759716), np.float64(0.07016183069707572), np.float64(0.0699987404471299), np.float64(0.06983816920329523), np.float64(0.0696800747956902), np.float64(0.06952441576676843), np.float64(0.06937115135926715), np.float64(0.06922024150437375), np.float64(0.06907164681008185), np.float64(0.06892532854974835), np.float64(0.0687812486508435), np.float64(0.0686393696838095), np.float64(0.06849965485159509), np.float64(0.06836206797815195), np.float64(0.06822657349874123), np.float64(0.06809313644919561), np.float64(0.067961722455845), np.float64(0.06783229772553254), np.float64(0.06770428903579932), np.float64(0.06757928372523506), np.float64(0.06745562968399212), np.float64(0.06733383534445969), np.float64(0.06721386967209597), np.float64(0.06709570215641501), np.float64(0.06697930280212627), np.float64(0.06686464212042395), np.float64(0.06675169112042348), np.float64(0.06664040213007429), np.float64(0.06653080846122665), np.float64(0.06642281359480932), np.float64(0.06631642107951677), np.float64(0.066211608074060279), np.float64(0.06610832559281864), np.float64(0.06600657021281309), np.float64(0.0659063105316164), np.float64(0.06580752017752023), np.float64(0.065710171519840698), np.float64(0.06561425155510119), np.float64(0.06551972561416586), np.float64(0.06542657414108709), np.float64(0.0653477429352925), np.float64(0.06524430361470467), np.float64(0.06515514002685512), np.float64(0.06506726182484374), np.float64(0.06498064766985515), np.float64(0.0648952765832028), np.float64(0.06481112794023773), np.float64(0.06472818146436811), np.float64(0.06464464172211699), np.float64(0.06456581561260431), np.float64(0.06448635737133043), np.float64(0.0644080235551142), np.float64(0.06433079554133217), np.float64(0.06425465582156798), np.float64(0.0641795839630046), np.float64(0.06410556476968135), np.float64(0.0640325799440141), np.float64(0.063960612416643), np.float64(0.06388964537111992), np.float64(0.06381966227619645), np.float64(0.06375064687909507), np.float64(0.06368258320118077), np.float64(0.0636154555332655), np.float64(0.06354924843135755), np.float64(0.06348394671157162), np.float64(0.06341953544633615), np.float64(0.063359999575896), np.float64(0.06329332582343267), np.float64(0.06323149885225086), np.float64(0.06317050510029515), np.float64(0.06311033085679153), np.float64(0.06305096264213547), np.float64(0.06299238720398384), np.float64(0.0629345915134133), np.float64(0.06287756276114324), np.float64(0.0628128835382297), np.float64(0.0627657559103815), np.float64(0.06271095325843989), np.float64(0.06265686843077901), np.float64(0.06260348966188053), np.float64(0.06255080538450809), np.float64(0.06249880422636036), np.float64(0.06244747500677472), np.float64(0.06239680673348793), np.float64(0.0623467859945137), np.float64(0.06229740997970036), np.float64(0.0622486604282762), np.float64(0.06220052967520031), np.float64(0.062153007623499706), np.float64(0.062106084346282515), np.float64(0.062059750083863094), np.float64(0.06201399524093575), np.float64(0.06196881038379625), np.float64(0.061924186237610215), np.float64(0.06188011368372787), np.float64(0.0618365837570441), np.float64(0.0618365837570441)
```

## To - Do - 8:

```
[14]: def rmse(y, y_pred):
    """
    Calculates the Root Mean Squared Error (RMSE) between actual and predicted values.

    Arguments:
    y: array-like
        Array of actual (target) values.
    y_pred: array-like
        Array of predicted values.

    Returns:
    float
        The root mean squared error.
    """
    Y = np.array(y, dtype = float).flatten()
    Y_pred = np.array(y_pred, dtype = float).flatten()

    rmse = np.sqrt(np.mean((Y - Y_pred) ** 2))
    return rmse
```

## To - Do - 9:

```
[15]: def r2(Y, Y_pred):
    """
    This function calculates the R Squared Error.

    Arguments:
    Y: array-like
        Array of actual (target) dependent values.
    Y_pred: array-like
        Array of predicted dependent values.

    Returns:
    float
        R Squared error.
    """
    Y = np.array(Y, dtype=float).flatten()
    Y_pred = np.array(Y_pred, dtype=float).flatten()

    mean_y = np.mean(Y) # Mean of actual values

    # Total sum of squares
    ss_tot = np.sum((Y - mean_y) ** 2)

    # Sum of squared residuals
    ss_res = np.sum((Y - Y_pred) ** 2)

    # R squared
    r2_score = 1 - (ss_res / ss_tot)

    return r2_score
```

## To - Do - 10:

```
[18]: # Step 1: Load the dataset
data = pd.read_csv('student.csv')

# Step 2: Split the data into features (X) and target (Y)
X = data[['Math', 'Reading']].values # Features: Math and Reading marks
Y = data['Writing'].values # Target: Writing marks

# Step 3: Split the data into training and test sets (80% train, 20% test)
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)

# Step 4: Initialize weights (W) to zeros, Learning rate and number of iterations
W = np.zeros(X_train.shape[1]) # Initialize weights
alpha = 0.00001 # Learning rate
iterations = 1000 # Number of iterations for gradient descent

# Step 5: Perform Gradient Descent
W_optimal, cost_history = gradient_descent(X_train, Y_train, W, alpha, iterations)

# Step 6: Make predictions on the test set
Y_pred = np.dot(X_test, W_optimal)

# Step 7: Evaluate the model using RMSE and R-Squared
model_rmse = rmse(Y_test, Y_pred)
model_r2 = r2(Y_test, Y_pred)

# Step 8: Output the results
print("Final Weights:", W_optimal)
print("Cost History (First 10 iterations):", cost_history[:10])
print("RMSE on Test Set:", model_rmse)
print("R-Squared on Test Set:", model_r2)

Final Weights: [[0.34811659]
 [0.64614558]]
Cost History (First 10 iterations): [np.float64(2013.165570783755), np.float64(1640.286832599692), np.float64(1337.0619994901588), np.float64(1090.4794892850578), np.float64(889.9583270083234), np.float64(726.8940993009545), np.float64(594.2897260808594), np.float64(486.4552052951635), np.float64(398.7634463599484), np.float64(327.4517147324688)]
RMSE on Test Set: 5.2798239764188635
R-Squared on Test Set: 0.8886354462786421
```

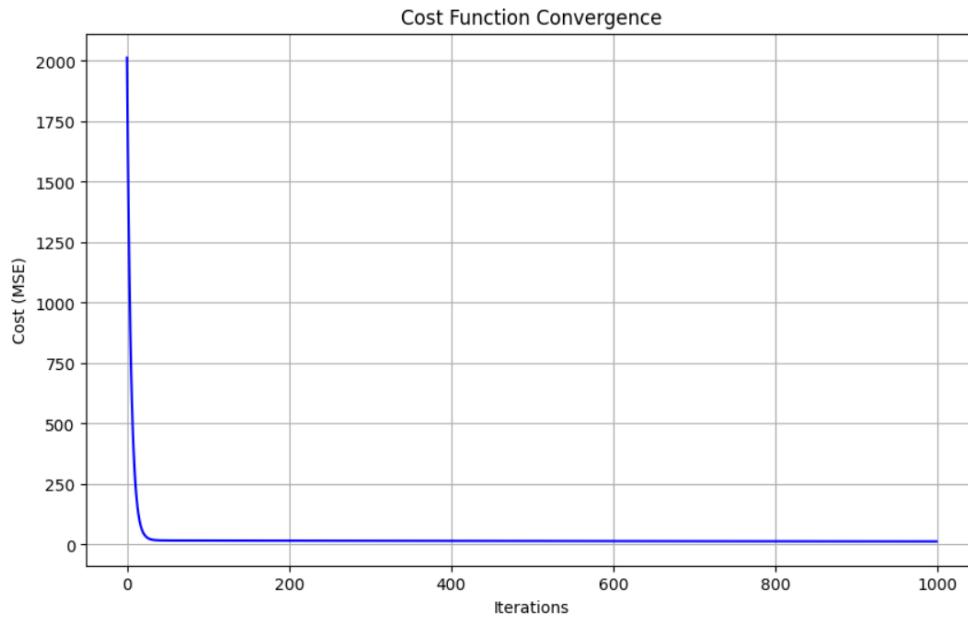
## To - Do - 11:

1. Did your Model Overfitt, Underfitts, or performance is acceptable.

```
import matplotlib.pyplot as plt

def plot_cost(cost_history):
    plt.figure(figsize=(10, 6))
    plt.plot(range(len(cost_history)), cost_history, color='blue')
    plt.title('Cost Function Convergence')
    plt.xlabel('Iterations')
    plt.ylabel('Cost (MSE)')
    plt.grid(True)
    plt.show()

plot_cost(cost_history)
```



The model's performance is acceptable and the cost function decreases smoothly indicating proper convergence. The RMSE value being around 5.28 shows that the prediction error is low and R-squared value of approx 0.89 indicates that the model explains most of the variance in writing marks. Therefore, the model neither underfits nor overfits the data.

## 2. Experiment with different value of learning rate, making it higher and lower, observe the result.

```
learning_rates = [0.000001, 0.00001, 0.0001]

for alpha in learning_rates:
    print("\nLearning rate:", alpha)

    W = np.zeros(X_train.shape[1])
    iterations = 1000

    W_optimal, cost_history = gradient_descent(
        X_train, Y_train, W, alpha, iterations
    )

    Y_pred = X_test @ W_optimal

    print(" Final Cost:", cost_history[-1])
    print(" RMSE:", rmse(Y_test, Y_pred))
    print(" R^2:", r2(Y_test, Y_pred))
```

```
Learning rate: 1e-06
Final Cost: 16.535602355147176
RMSE: 5.856694748793876
R^2: 0.8629707528684534
```

```
Learning rate: 1e-05
Final Cost: 13.150619992105618
RMSE: 5.2798239764188635
R^2: 0.8886354462786421
```

```
Learning rate: 0.0001
Final Cost: 10.26076310841341
RMSE: 4.792607360540954
R^2: 0.908240340333986
```

According to the data above, as the learning rate increases from 0.000001 to 0.0001, the model shows improved convergence with decrease in final cost and RMSE and increase in R-Squared. Therefore, smaller learning rate results in slower convergence while moderate learning rate results in quicker convergence and is more effective but with bigger learning rate it might start to diverge.