

```
[32]: # Rachit Manandhar
# 2501387
# np03cs4a240053

[33]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import sklearn

from sklearn.linear_model import Ridge, Lasso
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.metrics import mean_squared_error, accuracy_score
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
```

3. Regression Task (California Housing)

Task 1: Load and Split Dataset

Loading the Dataset:

```
[34]: # Setting the columns
cols = ["longitude", "latitude", "housingMedianAge", "totalRooms", "totalBedrooms", "population", "households", "medianIncome", "medianHouseValue"]

# Loading the dataset by downloading from "https://s3-eu-west-1.amazonaws.com/pfigshare-u-files/5976036/cal_housing.tgz"
df = pd.read_csv("cal_housing.data", header=None, names=cols)

# Splitting the data into features and label
X = df.drop("medianHouseValue", axis = 1)
y = df["medianHouseValue"]

# Train-test split (80-20)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print("Training samples:", X_train.shape)
print("Test samples:", X_test.shape)

Training samples: (16512, 8)
Test samples: (4128, 8)
```

Task 2:

2.1:

```
[35]: lr = LinearRegression()
lr.fit(X_train, y_train)

y_train_pred = lr.predict(X_train)
y_test_pred = lr.predict(X_test)

lr_train_mse = mean_squared_error(y_train, y_train_pred)
lr_test_mse = mean_squared_error(y_test, y_test_pred)

print("Baseline Linear Regression")
print("Train MSE:", lr_train_mse)
print("Test MSE:", lr_test_mse)

print("Coefficients:", lr.coef_)

Baseline Linear Regression
Train MSE: 4811134397.884198
Test MSE: 4918556441.477801
Coefficients: [-4.26323917e+04 -4.24500719e+04  1.18280965e+03 -8.18797708e+00
 1.16260128e+02 -3.84922131e+01  4.63425720e+01  4.05384044e+04]
```

2.2:

```
[36]: # Ridge hyperparameter tuning
ridge = Ridge()

alpha_grid = {'alpha': [0.001, 0.01, 0.1, 1, 10, 100]}

ridge_cv = GridSearchCV(ridge, alpha_grid, scoring='neg_mean_squared_error', cv = 5)

ridge_cv.fit(X_train, y_train)

print("Best Ridge alpha:", ridge_cv.best_params_)

Best Ridge alpha: {'alpha': 10}

[37]: best_ridge = ridge_cv.best_estimator_

# Evaluate Ridge
y_train_pred = best_ridge.predict(X_train)
y_test_pred = best_ridge.predict(X_test)

ridge_train_mse = mean_squared_error(y_train, y_train_pred)
ridge_test_mse = mean_squared_error(y_test, y_test_pred)
print("Ridge Train MSE:", ridge_train_mse)
print("Ridge Test MSE:", ridge_test_mse)

Ridge Train MSE: 4811139082.000713
Ridge Test MSE: 4918567284.46597

[38]: lasso = Lasso(max_iter=10000)

lasso_cv = GridSearchCV(lasso, alpha_grid, scoring='neg_mean_squared_error', cv = 5)

lasso_cv.fit(X_train, y_train)

print("Best Lasso alpha:", lasso_cv.best_params_)

Best Lasso alpha: {'alpha': 10}

[39]: best_lasso = lasso_cv.best_estimator_

# Evaluate Lasso
y_train_pred = best_lasso.predict(X_train)
y_test_pred = best_lasso.predict(X_test)

lasso_train_mse = mean_squared_error(y_train, y_train_pred)
lasso_test_mse = mean_squared_error(y_test, y_test_pred)
print("Lasso Train MSE:", lasso_train_mse)
print("Lasso Test MSE:", lasso_test_mse)

Lasso Train MSE: 4811135093.259237
Lasso Test MSE: 4918555581.562368
```

2.3:

```
[40]: coef_comparison = pd.DataFrame({
    "Feature": X.columns,
    "Ridge Coef": best_ridge.coef_,
    "Lasso Coef": best_lasso.coef_
})

print(coef_comparison)


print("\nPerformance Comparison:")
print(f"{'Model':<10} {'Train MSE':<15} {'Test MSE':<15}")
print()
print(f"{'Baseline':<10} {lr_train_mse:<15f} {lr_test_mse:<15f}")
print(f"{'Ridge':<10} {ridge_train_mse:<15f} {ridge_test_mse:<15f}")
print(f"{'Lasso':<10} {lasso_train_mse:<15f} {lasso_test_mse:<15f}")

plt.figure(figsize=(12,5))
plt.plot(best_ridge.coef_, label="Ridge (L2)")
plt.plot(best_lasso.coef_, label="Lasso (L1)")
plt.legend()
plt.title("Coefficient Comparison: Ridge vs Lasso")
plt.show()
```

	Feature	Ridge Coef	Lasso Coef
0	longitude	-42535.627082	-42595.282794
1	latitude	-42359.666504	-42415.402048
2	housingMedianAge	1184.351988	1183.328980
3	totalRooms	-8.196936	-8.191118
4	totalBedrooms	116.124492	116.204998
5	population	-38.496151	-38.493803
6	households	46.569026	46.430671
7	medianIncome	40543.565513	40540.088671

Performance Comparison:

Model	Train MSE	Test MSE
Baseline	4811134397.884198	4918556441.477801
Ridge	4811139082.000713	4918567284.465970
Lasso	4811135093.259237	4918555581.562368



4. Classification Task (Breast Cancer)

▼ Task 1: Load and Split Dataset

```
[41]: from sklearn.datasets import load_breast_cancer

[42]: X, y = load_breast_cancer(return_X_y=True)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)

print("Training samples:", X_train.shape)
print("Test samples:", X_test.shape)

Training samples: (455, 30)
Test samples: (114, 30)
```

▼ Task 2:

1.

```
[43]: # Baseline Logistic Regression
log_reg = LogisticRegression(max_iter=10000)
log_reg.fit(X_train, y_train)

# Predictions
y_train_pred = log_reg.predict(X_train)
y_test_pred = log_reg.predict(X_test)

# Accuracy
print("Baseline Logistic Regression")
print("Train Accuracy:", accuracy_score(y_train, y_train_pred))
print("Test Accuracy:", accuracy_score(y_test, y_test_pred))

# Coefficients
print("Coefficients:", log_reg.coef_)

Baseline Logistic Regression
Train Accuracy: 0.9626373626373627
Test Accuracy: 0.956140350877193
Coefficients: [[ 0.98208299  0.22519686 -0.36688444  0.0262268  -0.15507824 -0.22867976
 -0.52338614 -0.2793554  -0.22391176 -0.03605388 -0.09476544  1.39135347
 -0.16429246 -0.08903006 -0.02250974  0.04944847 -0.04186075 -0.03193634
 -0.03298528  0.01189208  0.10400464 -0.51389384 -0.01711567 -0.01662253
 -0.30695364 -0.75341491 -1.41533107 -0.50382259 -0.73542849 -0.09913574]]
```

2.

```
[56]: param_grid = {'C': [0.001, 0.01, 0.1, 1, 10, 100], 'solver': ['liblinear']}
log_cv = GridSearchCV(log_reg, param_grid, scoring = 'accuracy', cv = 5)
log_cv.fit(X_train, y_train)

print("Best Parameters:", log_cv.best_params_)
```

Best Parameters: {'C': 10, 'solver': 'liblinear'}

```
[49]: best_log = log_cv.best_estimator_

y_train_pred = best_log.predict(X_train)
y_test_pred = best_log.predict(X_test)

print("Tuned Logistic Regression")
print("Train Accuracy:", accuracy_score(y_train, y_train_pred))
print("Test Accuracy:", accuracy_score(y_test, y_test_pred))

Tuned Logistic Regression
Train Accuracy: 0.9692307692307692
Test Accuracy: 0.956140350877193
```

```
[59]: # L1 Logistic Regression
log_l1 = LogisticRegression(C=log_cv.best_params_['C'], l1_ratio = 1.0, solver='liblinear', max_iter=10000)
log_l1.fit(X_train, y_train)

# L2 Logistic Regression
log_l2 = LogisticRegression(C=log_cv.best_params_['C'], l1_ratio = 0.0, solver='liblinear', max_iter=10000)
log_l2.fit(X_train, y_train)

print("L1 Train Accuracy:", accuracy_score(y_train, log_l1.predict(X_train)))
print("L1 Test Accuracy:", accuracy_score(y_test, log_l1.predict(X_test)))

print("L2 Train Accuracy:", accuracy_score(y_train, log_l2.predict(X_train)))
print("L2 Test Accuracy:", accuracy_score(y_test, log_l2.predict(X_test)))

L1 Train Accuracy: 0.9824175824175824
L1 Test Accuracy: 0.9736842105263158
L2 Train Accuracy: 0.9692307692307692
L2 Test Accuracy: 0.956140350877193
```

3.

```
[60]: # Coefficients
coeff_l1 = pd.Series(log_l1.coef_[0], index=load_breast_cancer().feature_names)
coeff_l2 = pd.Series(log_l2.coef_[0], index=load_breast_cancer().feature_names)

print("L1 coefficients:")
print(coeff_l1)

print("\nL2 coefficients:")
print(coeff_l2)

# Count number of non-zero coefficients
print("\nNon-zero L1 coefficients:", np.sum(coeff_l1 != 0))
print("Non-zero L2 coefficients:", np.sum(coeff_l2 != 0))

L1 coefficients:
mean radius           1.643271
mean texture          0.179620
mean perimeter        -0.051568
mean area              -0.011643
mean smoothness        0.000000
mean compactness       0.000000
mean concavity         0.000000
mean concave points   -14.210292
mean symmetry          0.000000
mean fractal dimension 0.000000
radius error           0.000000
texture error          3.221484
perimeter error        -0.874976
area error              -0.091484
smoothness error       0.000000
compactness error      0.000000
concavity error        2.468309
concave points error   0.000000
symmetry error         0.000000
fractal dimension error 0.000000
worst radius            0.871757
worst texture           -0.600484
worst perimeter         0.135730
worst area              -0.026988
worst smoothness        0.000000
worst compactness       0.000000
worst concavity         -2.272920
worst concave points   -31.216774
worst symmetry           -6.645943
worst fractal dimension 0.000000
dtype: float64
```

```
L2 coefficients:  
mean radius           4.398356  
mean texture          0.296195  
mean perimeter        -0.498847  
mean area             -0.008948  
mean smoothness       -0.546015  
mean compactness      -0.819821  
mean concavity        -1.604845  
mean concave points   -1.233964  
mean symmetry         -0.739526  
mean fractal dimension -0.022000  
radius error          -0.321500  
texture error          3.604604  
perimeter error       -0.969272  
area error            -0.071975  
smoothness error      -0.074886  
compactness error     0.424390  
concavity error       0.441039  
concave points error  -0.100238  
symmetry error        -0.080574  
fractal dimension error 0.084301  
worst radius          0.398100  
worst texture          -0.672205  
worst perimeter        0.188310  
worst area             -0.027294  
worst smoothness       -1.040566  
worst compactness      -2.185084  
worst concavity        -3.241211  
worst concave points   -2.042643  
worst symmetry         -2.605688  
worst fractal dimension -0.153856  
dtype: float64
```

Non-zero L1 coefficients: 16

Non-zero L2 coefficients: 30

```
[61]: print("Accuracy Comparison: ")
results = pd.DataFrame({
    'Model': ['L1', 'L2'],
    'Train Accuracy': [
        accuracy_score(y_train, log_l1.predict(X_train)),
        accuracy_score(y_train, log_l2.predict(X_train))
    ],
    'Test Accuracy': [
        accuracy_score(y_test, log_l1.predict(X_test)),
        accuracy_score(y_test, log_l2.predict(X_test))
    ]
})
print(results)
```

```
Accuracy Comparison:
   Model  Train Accuracy  Test Accuracy
0      L1      0.982418      0.973684
1      L2      0.969231      0.956140
```

Discussion:

1. L1 produces sparse coefficients (feature selection)
2. L2 shrinks all coefficients but rarely zero
3. Regularization reduces variance and mitigates overfitting
4. Overly strong regularization may increase bias

Optional:

```
[62]: plt.figure(figsize=(12,5))

plt.plot(log_l1.coef_[0], label="L1")
plt.plot(log_l2.coef_[0], label="L2")

plt.legend()
plt.title("Logistic Regression Coefficients (L1 vs L2)")
plt.show()
```

